

# SKYTRAX AIRLINE REVIEW ANALYSIS

Rounak Saha

2024-12-17

```
options(repos = c(CRAN = "https://cloud.r-project.org"))
install.packages("tidyverse")
```

```
##
## The downloaded binary packages are in
## /var/folders/fm/k4r2n81j77q8l80mh7tnvq2r0000gn/T//RtmpXdsoRR/downloaded_packages
```

```
install.packages("rvest")
```

```
##
## The downloaded binary packages are in
## /var/folders/fm/k4r2n81j77q8l80mh7tnvq2r0000gn/T//RtmpXdsoRR/downloaded_packages
```

```
install.packages("data.table")
```

```
##
## The downloaded binary packages are in
## /var/folders/fm/k4r2n81j77q8l80mh7tnvq2r0000gn/T//RtmpXdsoRR/downloaded_packages
```

```
install.packages("wordcloud")
```

```
##
## The downloaded binary packages are in
## /var/folders/fm/k4r2n81j77q8l80mh7tnvq2r0000gn/T//RtmpXdsoRR/downloaded_packages
```

```
library(tidyverse)
```

```
## -- Attaching core tidyverse packages ----- tidyverse 2.0.0 --
## v dplyr      1.1.4      v readr      2.1.5
## v forcats    1.0.0      v stringr   1.5.0
## v ggplot2    3.5.1      v tibble    3.2.1
## v lubridate  1.9.3      v tidyr     1.3.1
## v purrr      1.0.2
```

```
## -- Conflicts ----- tidyverse_conflicts() --
```

```
## x dplyr::filter() masks stats::filter()
```

```
## x dplyr::lag()     masks stats::lag()
```

```
## i Use the conflicted package (<http://conflicted.r-lib.org/>) to force all conflicts to become errors
```

```
library(rvest)
```

```
##  
## Attaching package: 'rvest'  
##  
## The following object is masked from 'package:readr':  
##  
##     guess_encoding
```

```
library(data.table)
```

```
## Warning: package 'data.table' was built under R version 4.3.3
```

```
##  
## Attaching package: 'data.table'  
##  
## The following objects are masked from 'package:lubridate':  
##  
##     hour, isoweek, mday, minute, month, quarter, second, wday, week,  
##     yday, year  
##  
## The following objects are masked from 'package:dplyr':  
##  
##     between, first, last  
##  
## The following object is masked from 'package:purrr':  
##  
##     transpose
```

```
library(wordcloud)
```

```
## Loading required package: RColorBrewer
```

```
dataloop <- data.frame()
```

```
# for loop  
for (i in 1:30) {  
  
  # construct the URL with the current value of start  
  urlloop <- paste0("https://www.airlinequality.com/airline-reviews/british-airways/page/",  
                    i, "?sortby=post_date%3ADesc&pagesize=100")  
  
  # Read the html content of the webpage  
  webpageloop <- read_html(urlloop)  
  
  #Extract the data from the webpage using CSS selectors  
  Review = webpageloop %>% html_nodes(".text_content") %>% html_text()  
  # Extract all recommendations from multiple reviews  
  Recommended <- webpageloop %>%  
    html_nodes(".recommended+ .review-value") %>% # Select review recommendation nodes  
    html_text(trim = TRUE)
```

```

Date_Flowen = webpageloop %>% html_nodes(".date_flowen+ .review-value") %>% html_text()
Seat_Type =webpageloop %>% html_nodes(".cabin_flowen+ .review-value") %>% html_text()
Rating = webpageloop %>% html_nodes(".position-content .rating-10 span:nth-child(1)") %>% html_text()

#combine extracted data into a data table
data <- data.frame(Review,Date_Flowen,Rating,Recommended,Seat_Type)

data <- data %>%
  mutate(Review = ifelse(grepl("\\\\|", Review), Review, paste("\\", Review))) %>% # Add a placeholder
  separate(Review, into = c("Status", "Details"), sep = "\\|", fill = "right", extra = "merge") %>%
  mutate(Status = trimws(Status), Details = trimws(Details)) # Trim whitespace

# add the data to overall data table
dataloop <- bind_rows(dataloop,data)

#Pause for few secs to avoid overload
Sys.sleep(3)
}

view(dataloop)

```

```
str(dataloop)
```

```

## 'data.frame':   3000 obs. of  6 variables:
## $ Status      : chr  " Trip Verified" " Trip Verified" " Trip Verified" " Trip Verified" ...
## $ Details     : chr  "On a recent flight from Cyprus BA621 on 23/11/24, the second the cabin door wa
## $ Date_Flowen : chr  "November 2024" "December 2024" "December 2024" "December 2024" ...
## $ Rating      : chr  "1" "1" "2" "8" ...
## $ Recommended: chr  "no" "no" "no" "yes" ...
## $ Seat_Type   : chr  "Economy Class" "Economy Class" "Business Class" "Business Class" ...

```

```

dataloop$Rating <- as.numeric(dataloop$Rating)
dataloop <- as.data.table(dataloop)

```

```
dataloop[, .N, Rating]
```

```

##      Rating      N
##      <num> <int>
## 1:       1   830
## 2:       2   366
## 3:       8   279
## 4:       7   216
## 5:       4   202
## 6:      10   218
## 7:       3   354
## 8:       6   144
## 9:       9   209
## 10:      5   182

```

```
dataloop[, .N, Recommended]
```

```
##      Recommended      N
##      <char> <int>
## 1:      no 1951
## 2:      yes 1049
```

```
dataloop[, .(Average_rating= mean(Rating)), by= Seat_Type]
```

```
##      Seat_Type Average_rating
##      <char>      <num>
## 1: Economy Class      3.924066
## 2: Business Class     4.792657
## 3: First Class        5.694805
## 4: Premium Economy     4.487805
```

*#We will now convert this into a tibble. One column, called line, will have the line number (from 1 to*

```
Review_text <- tibble(review_no = 1: length(dataloop$Details), text= dataloop$Details)
Review_text
```

```
## # A tibble: 3,000 x 2
##   review_no text
##   <int> <chr>
## 1      1 "On a recent flight from Cyprus BA621 on 23/11/24, the second the ~
## 2      2 "Flight BA 0560 arrived in Rome on 11 December where ALL passenger~
## 3      3 "This was the first time I flew British Airways, and it was a huge~
## 4      4 "Pretty good flight but still some small things that can be improv~
## 5      5 "Check in was fine, but no priority/fast track lines for security ~
## 6      6 "British Airways is absolute rubbish. I had to fly to Amsterdam fo~
## 7      7 "The flight time was changed at the last minute without warning an~
## 8      8 " I'm so fraustrated. My flight was cancelled last minute, which~
## 9      9 "We have sat on this plane for an hour and forty five minutes awai~
## 10    10 "British Airways stranding my wife and I at Heathrow Airport for 2~
## # i 2,990 more rows
```

*#We can now tokenise the text using unnest\_tokens again, and then count the numbers*  
`library(tidytext)`

```
tokens<- Review_text %>% unnest_tokens(word, text)
tokens
```

```
## # A tibble: 499,910 x 2
##   review_no word
##   <int> <chr>
## 1      1 on
## 2      1 a
## 3      1 recent
## 4      1 flight
## 5      1 from
## 6      1 cyprus
```

```
## 7      1 ba621
## 8      1 on
## 9      1 23
## 10     1 11
## # i 499,900 more rows
```

```
#As we are only interested in words that are text, let's remove all words with
#digits and special characters from our set of words. We can do this using grepl().
tokens <- as.data.table(tokens)
tokens <- tokens[grepl("\\d", word)==FALSE, ]
tokens <- tokens[grepl("[:alpha:]", word), ]
tokens
```

```
##      review_no      word
##      <int>    <char>
## 1:      1      a
## 2:      1  flight
## 3:      1  cyprus
## 4:      1    the
## 5:      1    the
## ---
## 300598:    3000    home
## 300599:    3000      at
## 300600:    3000 midnight
## 300601:    3000     was
## 300602:    3000 starving
```

```
#Now that the text is tidy, it is easy to count the number of occurrences of each word. We can do this
tokens[, .N, word]
```

```
##      word      N
##      <char> <int>
## 1:      a 11385
## 2:  flight  5766
## 3:  cyprus    6
## 4:    the 25110
## 5:   cabin 1527
## ---
## 9185: appalling    1
## 9186:  arabia    1
## 9187:    ruh     2
## 9188: degrade    1
## 9189: cheddar    1
```

```
#Analysis is made difficult because the text contains lots of short words, like "the", "of" and "and",
data(stop_words)
stop_words
```

```
## # A tibble: 1,149 x 2
##   word      lexicon
```

```
##      <chr>      <chr>
## 1 a           SMART
## 2 a's         SMART
## 3 able        SMART
## 4 about       SMART
## 5 above       SMART
## 6 according   SMART
## 7 accordingly SMART
## 8 across      SMART
## 9 actually    SMART
## 10 after      SMART
## # i 1,139 more rows
```

*#remove certain words that will not be used to determine the positive or negative sentiment*

*# Custom stopwords*

```
custom_stopwords <- tibble(word = c('flight', 'ba', 'passenger', 'u', 'london',
                                     'airway', 'british', 'airline', 'heathrow',
                                     'plane', 'lhr', 'review', 'airways', 'flights'))
```

*# Combine with default stopwords*

```
updated_stopwords <- stop_words %>%
  bind_rows(custom_stopwords)
```

*# Check the updated stopwords*

```
print(updated_stopwords)
```

```
## # A tibble: 1,163 x 2
##   word      lexicon
##   <chr>     <chr>
## 1 a        SMART
## 2 a's      SMART
## 3 able     SMART
## 4 about    SMART
## 5 above    SMART
## 6 according SMART
## 7 accordingly SMART
## 8 across   SMART
## 9 actually SMART
## 10 after   SMART
## # i 1,153 more rows
```

*#We can remove these stop words from word column by performing an anti-join between tokens and stop\_words*

```
review_tokens <- tokens %>% anti_join(updated_stopwords)
```

## Joining with `by = join\_by(word)`

```
review_tokens
```


```
##      review_no    word
##      <int>     <char>
## 1:         1    cyprus
```

```
##      2:      1      cabin
##      3:      1      closed
##      4:      1      pilot
##      5:      1 announced
##      ---
## 123897:    3000    airport
## 123898:    3000    arrived
## 123899:    3000      home
## 123900:    3000  midnight
## 123901:    3000  starving
```

```
#Now, when we count the words, we will only get the meaningful words;
review_tokens[, .N, word][order(-N)]
```

```
##           word      N
##          <char> <int>
##      1:      seat 1836
##      2:      seats 1613
##      3:      class 1599
##      4:      cabin 1527
##      5:      staff 1266
##      ---
## 8761: undertaking      1
## 8762:    appalling      1
## 8763:      arabia      1
## 8764:      degrade      1
## 8765:      cheddar      1
```

```
review_tokens[, .N, word][order(-N)] %>% head(20) %>% ggplot(aes(N, reorder(word, N))) + geom_col(aes(f
  labs(x = "Count", y = "Word", title = "Top 20 Words") +
  theme_minimal()
```



```
# Sentiment Analysis
#The tidytext library provides the get_sentiments function. This can be used to download one of many di

sentiments<- get_sentiments("nrc")
sentiments
```

```
## # A tibble: 13,872 x 2
##   word      sentiment
##   <chr>     <chr>
## 1 abacus    trust
## 2 abandon   fear
## 3 abandon   negative
## 4 abandon   sadness
## 5 abandoned anger
## 6 abandoned fear
## 7 abandoned negative
## 8 abandoned sadness
## 9 abandonment anger
## 10 abandonment fear
## # i 13,862 more rows
```

```
# We can classify each word by sentiment by joining together the tidy text sherlock tibble with the dic

review_sentiments <- review_tokens %>% inner_join(sentiments)
```

```
## Joining with `by = join_by(word)`
```



```
## Warning in inner_join(., sentiments): Detected an unexpected many-to-many relationship between `x` and `y`.
## i Row 4 of `x` matches multiple rows in `y`.
## i Row 715 of `y` matches multiple rows in `x`.
## i If a many-to-many relationship is expected, set `relationship = "many-to-many"` to silence this warning.
```

```
review_sentiments
```

```
##      review_no      word      sentiment
##      <int>      <char>      <char>
##  1:          1      pilot      positive
##  2:          1      pilot      trust
##  3:          1      hell      anger
##  4:          1      hell      disgust
##  5:          1      hell      fear
##  ---
## 66679:      3000 appalling      negative
## 66680:      3000      hot      anger
## 66681:      3000      eat      positive
## 66682:      3000 airport anticipation
## 66683:      3000 starving      negative
```

```
#We can then count how many words of different sentiments there are using filter and count.
review_sentiments %>% filter(sentiment == "positive") %>% count(word, sort = TRUE)
```

```
##      word      n
##      <char> <int>
##  1:    flying  535
##  2:   friendly  462
##  3: entertainment  441
##  4:   breakfast  428
##  5:      pay    421
##  ---
## 726:   vocabulary      1
## 727:    warranty      1
## 728:   wholesome      1
## 729:    wireless      1
## 730: wonderfully      1
```

```
# We could get the total number of words of each type using;
review_sentiments %>% filter(sentiment == "positive") %>% count(word, sort = TRUE) %>% summarise(total = sum(n))
```


```
##      total
## 1 14670
```

```
#Alternatively (and much more simply) we could just count the number of occurrences of each sentiment using
review_sentiments %>% count(sentiment)
```

```
##      sentiment      n
##      <char> <int>
##  1:      anger  3558
```

```
## 2: anticipation 7034
## 3:    disgust  3391
## 4:    fear    4673
## 5:    joy     5351
## 6:   negative 11107
## 7:   positive 14670
## 8:    sadness  5594
## 9:    surprise 3121
## 10:   trust   8184
```

*#We could plot this using a similar technique as the last section (this time converting the sentiments*  
*review\_sentiments %>% count(sentiment) %>% ggplot(aes(n, reorder(sentiment, n))) + geom\_col(aes(fill= n*



*#What we want to do now is to count the number of words with different sentiments for different review*  
*review\_sentiments\_pivot<- review\_sentiments %>% count(review\_no,sentiment)*

*#Next, we can now pivot the tibble so that the sentiments are the columns, and we have one row per block*

```
review_sentiments_pivot <- review_sentiments_pivot %>% pivot_wider(names_from = sentiment, values_from = n)
review_sentiments_pivot
```


```
## # A tibble: 2,983 x 11
##   review_no anger anticipation disgust fear negative positive sadness surprise
##   <int> <int>      <int>    <int> <int>    <int>    <int>    <int>    <int>
## 1         1     4          3      3     4       7       2       5       1
```

```
## 2      2      2      5      2      2      5      7      4      3
## 3      3      0      3      1      0      1      6      1      0
## 4      4      3      4      0      4      9      17     3      1
## 5      5      2      4      3      4      6      5      3      0
## 6      6      3      2      4      3      4      3      4      0
## 7      7      1      0      1      2      2      2      1      0
## 8      8      4      2      4      4      5      2      2      1
## 9      9      3      0      2      2      4      1      2      1
## 10     10     2      1      2      5      4      3      3      0
## # i 2,973 more rows
## # i 2 more variables: trust <int>, joy <int>
```

```
#We can now do things like calculate the difference between the number of positive and negative sentiment
review_sentiments_pivot <- review_sentiments_pivot %>% mutate(sentiment = positive - negative)
review_sentiments_pivot
```

```
## # A tibble: 2,983 x 12
##   review_no anger anticipation disgust fear negative positive sadness surprise
##   <int> <int>      <int>      <int> <int>      <int>      <int>      <int>      <int>
## 1      1      4      3      3      4      7      2      5      1
## 2      2      2      5      2      2      5      7      4      3
## 3      3      0      3      1      0      1      6      1      0
## 4      4      3      4      0      4      9      17     3      1
## 5      5      2      4      3      4      6      5      3      0
## 6      6      3      2      4      3      4      3      4      0
## 7      7      1      0      1      2      2      2      1      0
## 8      8      4      2      4      4      5      2      2      1
## 9      9      3      0      2      2      4      1      2      1
## 10     10     2      1      2      5      4      3      3      0
## # i 2,973 more rows
## # i 3 more variables: trust <int>, joy <int>, sentiment <int>
```

```
#This could then be graphed, with the "sentiment" column on the y axis, and the review no on the x axis
# checking for first 300 reviews
review_sentiments_pivot %>% head(300) %>% ggplot(aes(review_no, sentiment)) + geom_col()
```



```
#Word clouds
#We are going to create word clouds from the text reviews
# The wordcloud function is very easy to use. You just need to pass in a column of words (review_tokens,
review_tokens_wordcloud <- review_tokens[, .N, word][order(-N)]
suppressWarnings(review_tokens_wordcloud %>% with(wordcloud(word,N, min.freq = 50,random.order=FALSE,col
```



#n-grams

*#It is very common that you want to tokenise using tokens that represent pairs of words. This is because  
#A token comprising n words is called an "n-gram" (or "ngram"). Tokenising on bigrams or n-grams enable*

```
Review_ngrams <- Review_text %>% unnest_tokens(ngram, text, token = "ngrams", n=3)
Review_ngrams
```

```
## # A tibble: 493,910 x 2
##   review_no ngram
##   <int> <chr>
## 1      1 1 on a recent
## 2      2 1 a recent flight
## 3      3 1 recent flight from
## 4      4 1 flight from cyprus
## 5      5 1 from cyprus ba621
## 6      6 1 cyprus ba621 on
## 7      7 1 ba621 on 23
## 8      8 1 on 23 11
## 9      9 1 23 11 24
## 10     10 1 11 24 the
## # i 493,900 more rows
```

*#In this case we have placed the ngrams into the column ngram.*

#Now we want to remove pairs of words where one in the pair is a stop word. To do this we must first se

```
Review_ngrams <- Review_ngrams %>% separate(ngram, c("word1", "word2", "word3"), sep = " ")
Review_ngrams
```

```
## # A tibble: 493,910 x 4
##   review_no word1 word2 word3
##   <int> <chr> <chr> <chr>
## 1         1 on a recent
## 2         1 a recent flight
## 3         1 recent flight from
## 4         1 flight from cyprus
## 5         1 from cyprus ba621
## 6         1 cyprus ba621 on
## 7         1 ba621 on 23
## 8         1 on 23 11
## 9         1 23 11 24
## 10        1 11 24 the
## # i 493,900 more rows
```

*#Next we need to filter only the rows where neither word1 or word2 or word3 have a word in stop\_words.*

```
Review_ngrams <- Review_ngrams %>%
  filter(!word1 %in% updated_stopwords$word) %>%
  filter(!word2 %in% updated_stopwords$word) %>%
  filter(!word3 %in% updated_stopwords$word)
```

```
Review_ngrams
```

```
## # A tibble: 17,687 x 4
##   review_no word1 word2 word3
##   <int> <chr> <chr> <chr>
## 1         1 23 11 24
## 2         1 6 hours due
## 3         1 2 hours waiting
## 4         1 8 hour hell
## 5         1 3 oatmeal biscuits
## 6         1 4 cabin crew
## 7         1 23 11 24
## 8         1 email system unsurprisingly
## 9         2 sunday 15 december
## 10        4 chosen wines disappointingly
## # i 17,677 more rows
```

*#Next, we can see that there are a lot of NAs in the tibble. These are caused by blank lines or other t*

```
Review_ngrams <- Review_ngrams %>%
  filter(!is.na(word1)) %>%
  filter(!is.na(word2)) %>%
  filter(!is.na(word3))
```

```
Review_ngrams
```

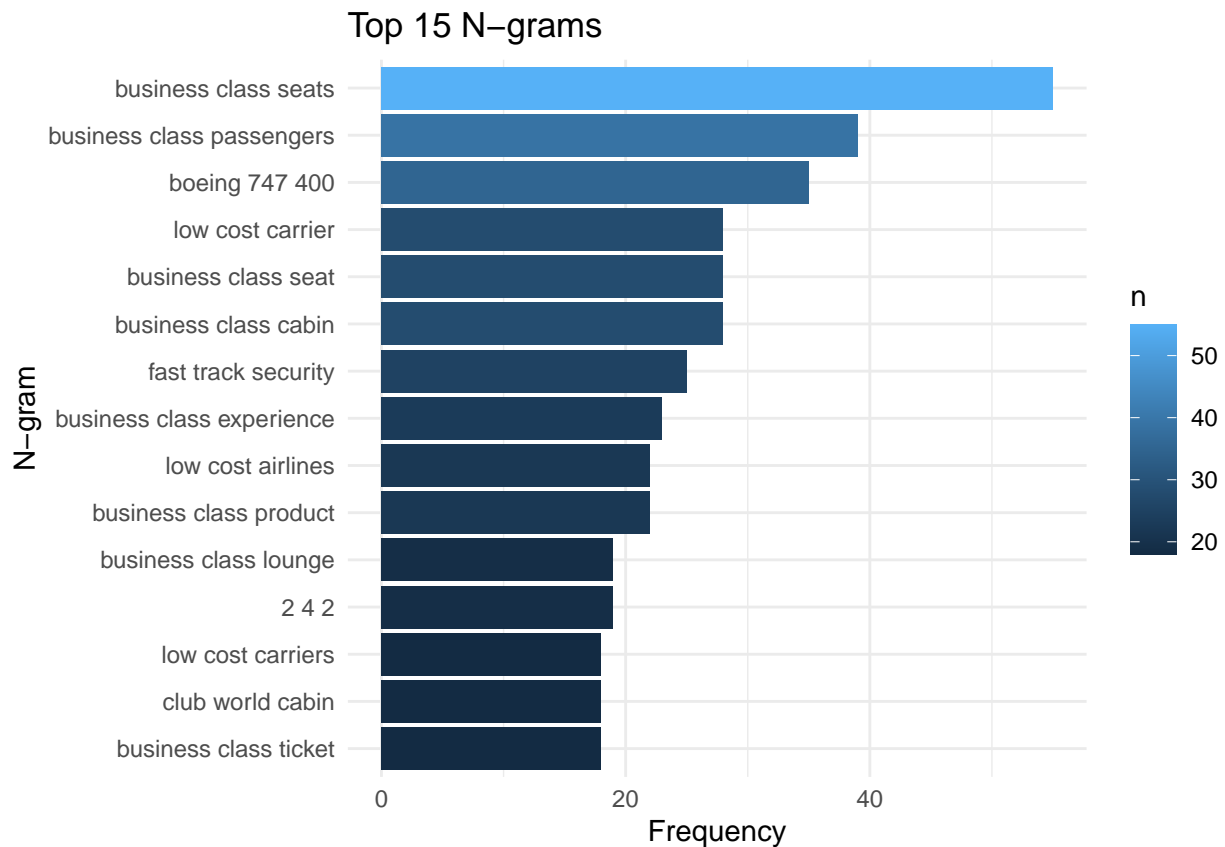
```
## # A tibble: 17,687 x 4
##   review_no word1 word2 word3
##   <int> <chr> <chr> <chr>
```

```
## 1      1 23      11      24
## 2      1 6       hours    due
## 3      1 2       hours    waiting
## 4      1 8       hour     hell
## 5      1 3       oatmeal  biscuits
## 6      1 4       cabin    crew
## 7      1 23      11      24
## 8      1 email   system   unsurprisingly
## 9      2 sunday  15      december
## 10     4 chosen  wines    disappointingly
## # i 17,677 more rows
```

```
#Finally! we will rejoin the word1 and word2 columns into a single "ngram" column, using the "unite" function
Review_ngrams <- Review_ngrams %>% unite(ngram, word1, word2, word3, sep = " ")
Review_ngrams
```

```
## # A tibble: 17,687 x 2
##   review_no ngram
##   <int> <chr>
## 1      1 1 23 11 24
## 2      2 1 6 hours due
## 3      3 1 2 hours waiting
## 4      4 1 8 hour hell
## 5      5 1 3 oatmeal biscuits
## 6      6 1 4 cabin crew
## 7      7 1 23 11 24
## 8      8 1 email system unsurprisingly
## 9      9 2 sunday 15 december
## 10    10 4 chosen wines disappointingly
## # i 17,677 more rows
```

```
#We can now count the ngrams as before, e.g.
ngram_counts <- Review_ngrams %>% count(ngram, sort=TRUE)
ngram_counts %>%
  head(15) %>%
  ggplot(aes(n, reorder(ngram, n))) + # Reorder ngram based on n in descending order
  geom_col(aes(fill= n)) +
  labs(x = "Frequency", y = "N-gram", title = "Top 15 N-grams") +
  theme_minimal()
```



*# Visualising correlations between words*

*#We can go beyond just counting, and can draw graphs that visualise the connections and correlations between words*

```
Review_ngrams_count <- Review_text %>% unnest_tokens(ngram, text, token = "ngrams", n=3) %>%
  separate(ngram, c("word1", "word2", "word3"), sep = " ") %>%
  filter(!word1 %in% updated_stopwords$word) %>%
  filter(!word2 %in% updated_stopwords$word) %>%
  filter(!word3 %in% updated_stopwords$word) %>%
  filter(!is.na(word1)) %>%
  filter(!is.na(word2)) %>%
  filter(!is.na(word3)) %>%
  count(word1, word2, word3, sort = TRUE)
```

Review\_ngrams\_count

```
## # A tibble: 15,604 x 4
##   word1 word2 word3      n
##   <chr> <chr> <chr>   <int>
## 1 business class seats    55
## 2 business class passengers 39
## 3 boeing 747 400        35
## 4 business class cabin    28
## 5 business class seat     28
## 6 low cost carrier        28
## 7 fast track security     25
```



```
## 8 business class experience    23
## 9 business class product      22
## 10 low      cost  airlines     22
## # i 15,594 more rows
```

*#This table can now be converted into a directed graph. A directed graph is one where nodes (in this case words) are connected by directed edges (in this case word pairs).*

*#To prevent the graph becoming too large, we will only graph pairs of words that appear seven or more times in the data.*

```
install.packages("igraph")
```

```
##
## The downloaded binary packages are in
## /var/folders/fm/k4r2n81j77q8l80mh7tnvq2r0000gn/T//RtmpXdsoRR/downloaded_packages
```

```
library(igraph)
```

```
## Warning: package 'igraph' was built under R version 4.3.3
```

```
##
## Attaching package: 'igraph'
```

```
## The following objects are masked from 'package:lubridate':
##
##    %--%, union
```

```
## The following objects are masked from 'package:dplyr':
##
##    as_data_frame, groups, union
```

```
## The following objects are masked from 'package:purrr':
##
##    compose, simplify
```

```
## The following object is masked from 'package:tidyr':
##
##    crossing
```

```
## The following object is masked from 'package:tibble':
##
##    as_data_frame
```

```
## The following objects are masked from 'package:stats':
##
##    decompose, spectrum
```

```
## The following object is masked from 'package:base':
##
##    union
```

```
ngram_graph <- Review_ngrams_count %>% filter(n > 5) %>%
  graph_from_data_frame()
ngram_graph
```

```
## IGRAPH 25e6669 DN-- 77 85 --
## + attr: name (v/c), word3 (e/c), n (e/n)
## + edges from 25e6669 (vertex names):
## [1] business->class    business->class    boeing    ->747    business->class
## [5] business->class    low    ->cost    fast    ->track    business->class
## [9] business->class    low    ->cost    2    ->4    business->class
## [13] business->class    club    ->world    low    ->cost    worst    ->business
## [17] premium ->economy    gold    ->card    business->class    club    ->world
## [21] haul    ->business    club    ->europe    premium ->economy    premium ->economy
## [25] 3    ->3    booked    ->business    cabin    ->crew    economy ->class
## [29] middle ->seat    poor    ->customer    boeing    ->777    business->class
## + ... omitted several edges
```

*#The graph\_from\_data\_frame takes a tibble where the first three columns name the nodes and specify their*

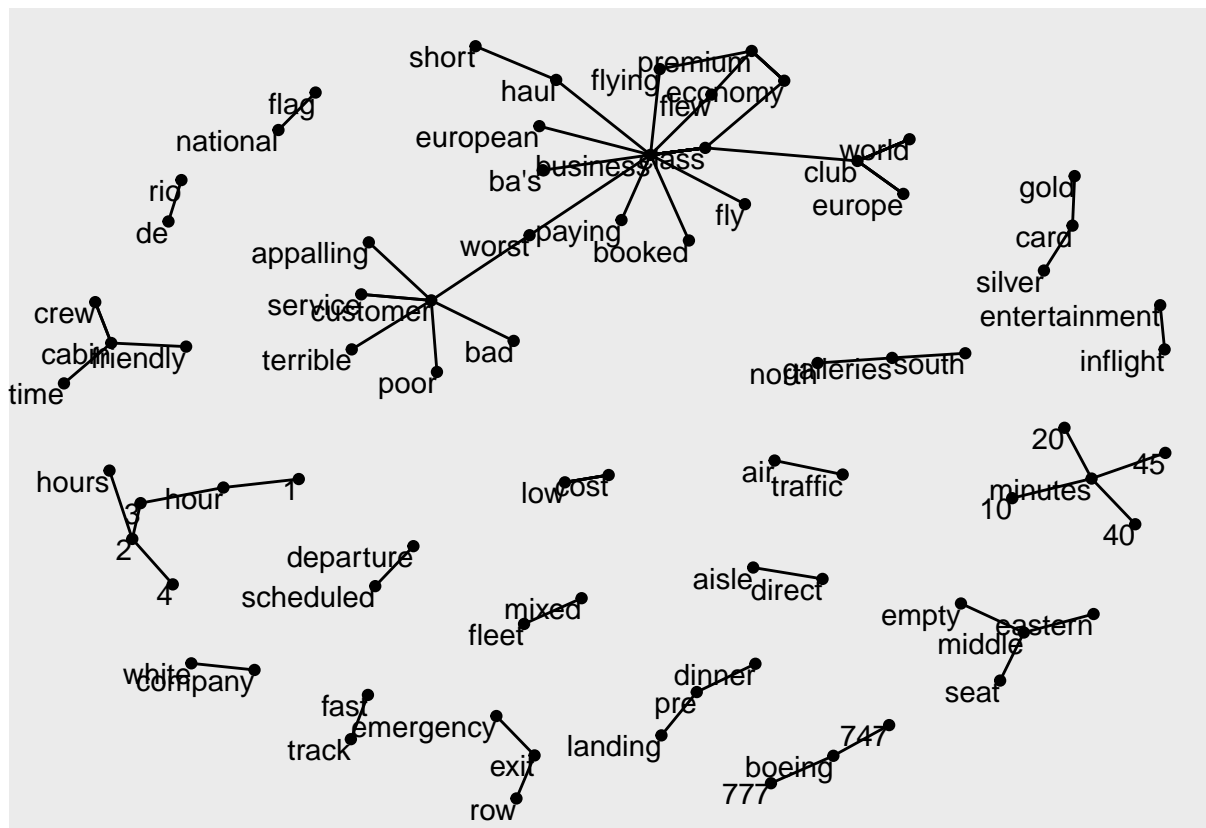
*#Now that we have built the directed graph, we can visualise it using the ggraph function from the ggraph*

*#There is a grammar for these graphs, e.g. we pass the data into ggraph, and then add layers, such as g*  
`install.packages("ggraph")`

```
##
## The downloaded binary packages are in
## /var/folders/fm/k4r2n81j77q8l80mh7tnvq2r0000gn/T//RtmpXdsoRR/downloaded_packages
```

```
library(ggraph)

ggraph(ngram_graph, layout = "fr") +
  geom_edge_link() +
  geom_node_point() +
  geom_node_text(aes(label = name), vjust = 1, hjust = 1)
```



*#This graph isn't very pretty. By looking online at the options to improve ggraph, and comparing against*

```
ggraph(ngram_graph, layout = "fr") +
  geom_edge_link(aes(edge_alpha = n), show.legend = FALSE,
    arrow = grid::arrow(type = "closed", length = unit(2, "mm")),
    end_cap = circle(1, "mm")) +
  geom_node_point(color = "lightblue", size = 2) +
  geom_node_text(aes(label = name), size = 2) +
  theme_void()
```

